

CS 332/532 – 1G- Systems Programming

Lab 11

Objectives

- Create threads using POSIX threads library
- Thread synchronization using Mutexes

Lab Assignment #11

Modify the pthread_sum.c program to create a structure and pass the structure as argument to the thread creation function instead of using global variables a, sum, N, and size. You have to create a structure that contains the variables a and sum with type double, variables N and size with type int, and variable tid with type long or int. You have to create an instance of this structure specific to each thread and pass the structure as an argument to the corresponding thread creation function. Test the program for different values of N and number of threads and make sure that the result is correct.

Assignment Submission

The assignment submission in Canvas is required. No late submissions will be accepted.

Submission Checklist:

- Upload the C source file (.c file) to Canvas as part of this lab submission.
Submissions through the Canvas “Comments” will not be accepted.
- Upload a README.md file which should include:
 - Instructions on how to compile your C source file into an executable.
 - How to run the executable program.
 - Any citation documentation.
 - A link to your GitHub repository.

Please do not upload executables or object files. Independent Completion Forms are not required for labs.

Lab Workbook

Create threads using POSIX threads library

In the previous labs we focused on how to create processes, in this lab we will focus on creating threads and mechanisms for establishing synchronization among threads.

First, let us understand the difference between a process and a thread. A process could be considered to have two characteristics: (a) resource ownership and (b) scheduling or execution. The unit of scheduling and dispatching is usually referred to as a thread or lightweight process and the ability of to support multiple, concurrent paths of execution within a single process is often referred to as *multithreading*. Threads offer several benefits compared to a process:

- Threads takes less time to create a new thread than a process
- Threads take less time to terminate a thread than a process
- Switching between two threads (context switching) takes less time than switching between processes
- All of the threads in a process share the state and resources of that process (since threads reside in the same address space and have access to the same data)
- Threads enhance efficiency in communication between programs (since threads share memory and files within the same process and can communicate without invoking the kernel)

As a result of the above advantages, if we have to implement a set of functions that are closely related, implementing this functionality using multiple threads is far more efficient than using multiple processes.

We will use the POSIX threads library, usually referred to as Pthreads library, that provides C APIs to create and manage threads. We have to include the file *pthread.h* and link with *-lpthread* to compile and link.

We can create new threads using the *pthread_create()* function which has the following function definition:

```
#include <pthread.h>
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,
                  void *(*start_routine) (void *), void *arg);
```

The new thread that will be created by the *pthread_create* function will invoke the function *start_routine*. Note that the function *start_routine* takes one argument of type *void*

* and has the return type as *void* *. In other words, the function *start_routine* has the following function definition:

```
void *start_routine(void *arg)
```

When the *pthread_create* call returns successfully, it returns the thread ID associated with the new thread created in the variable *thread*. This can be used by the main thread in subsequent *pthread* function calls such as *pthread_join*. The second argument, *attr*, provides a reference to the *pthread_attr_t* structure that describes the various attributes of the new thread to be created. It can be initialized using *pthread_attr_init* call or set to NULL if default attributes must be used. You can find out more about the different thread attributes that can be specified by looking at the man page for *pthread_attr_init*.

The new thread created will terminate when the function *start_routine* returns or when a call to *pthread_exit* is made inside the *start_routine*. We can use the *pthread_join* function to wait for a thread to complete using the thread ID that was returned when *pthread_create* call was invoked. If a thread has already completed, *pthread_join* will return immediately, otherwise, it will wait for the corresponding thread to complete.

The following example shows how to create threads and wait for the threads to complete. You can download this program here: [pthread1.c](#)

```
/*
Simple Pthread Program to illustrate the create/join threads.
To Compile: gcc -O -Wall pthread1.c -lpthread
To Run: ./a.out 4
*/
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

int nthreads; /* variable shared between threads */

void *compute(void *arg) {
    long tid = (long)arg;

    printf("Hello, I am thread %ld of %d\n", tid, nthreads);

    return (NULL);
}

int main(int argc, char **argv) {
    long i;
    pthread_t *tid;

    if (argc != 2) {
        printf("Usage: %s <# of threads>\n", argv[0]);
    }
}
```

```

    exit(-1);
}

nthreads = atoi(argv[1]); // no. of threads

// allocate vector and initialize
tid = (pthread_t *)malloc(sizeof(pthread_t)*nthreads);

// create threads
for ( i = 0; i < nthreads; i++)
    pthread_create(&tid[i], NULL, compute, (void *)i);

// wait for them to complete
for ( i = 0; i < nthreads; i++)
    pthread_join(tid[i], NULL);

printf("Exiting main program\n");

return 0;
}

```

Here is an updated version of the program that prints the thread ids: pthread2.c

Here is another version of the program that passes a structure as the argument for `pthread_create` instead of using the global variables: pthread3.c

Thread Synchronization using Mutexes

We can use the mutexes provided by the Pthreads library to control access to critical sections of the program and provide synchronization across the threads. We will use the example of computing the sum of the elements in a vector to illustrate the use of mutexes. We have a vector of N elements, we would like to assign each thread to compute the partial sum of N/P elements (where P is the number of threads), then we will update the shared global variable *sum* with the partial sums using mutex locks. The API for the mutex lock and unlock functions are shown below:

```

#include <pthread.h>
int pthread_mutex_lock(pthread_mutex_t *mutex);
int pthread_mutex_unlock(pthread_mutex_t *mutex);

```

The *mutex* variable is of type `pthread_mutex_t` and can be initially statically by assigning the value `PTHREAD_MUTEX_INITIALIZER`. Note that the mutex variable must be declared in global scope since it will be shared among multiple threads. A mutex can also be initialized dynamically using the function `pthread_mutex_init`. The `pthread_mutex_destroy` function can

be used to destroy the mutex that was initialized using `pthread_mutex_init`. The complete program is shown below and can be downloaded here: [pthread_sum.c](#)

```
/*
Simple Pthread Program to find the sum of a vector.
Uses mutex locks to update the global sum.
Author: Purushotham Bangalore
Date: Jan 25, 2009

To Compile: gcc -O -Wall pthread_sum.c -lpthread
To Run: ./a.out 1000 4
*/



#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <unistd.h>

pthread_mutex_t mutex=PTHREAD_MUTEX_INITIALIZER;

double *a=NULL, sum=0.0;
int N, size;

void *compute(void *arg) {
    int myStart, myEnd, myN, i;
    long tid = (long)arg;

    // determine start and end of computation for the current thread
    myN = N/size;
    myStart = tid*myN;
    myEnd = myStart + myN;
    if (tid == (size-1)) myEnd = N;

    // compute partial sum
    double mysum = 0.0;
    for (i=myStart; i<myEnd; i++)
        mysum += a[i];

    // grab the lock, update global sum, and release lock
    pthread_mutex_lock(&mutex);
    sum += mysum;
    pthread_mutex_unlock(&mutex);

    return (NULL);
}

int main(int argc, char **argv) {
    long i;
    pthread_t *tid;

    if (argc != 3) {
        printf("Usage: %s <# of elements> <# of threads>\n", argv[0]);
        exit(-1);
    }
```

```
}

N = atoi(argv[1]); // no. of elements
size = atoi(argv[2]); // no. of threads

// allocate vector and initialize
tid = (pthread_t *)malloc(sizeof(pthread_t)*size);
a = (double *)malloc(sizeof(double)*N);
for (i=0; i<N; i++)
    a[i] = (double)(i + 1);

// create threads
for ( i = 0; i < size; i++)
    pthread_create(&tid[i], NULL, compute, (void *)i);

// wait for them to complete
for ( i = 0; i < size; i++)
    pthread_join(tid[i], NULL);

printf("The total is %g, it should be equal to %g\n",
      sum, ((double)N*(N+1))/2);

return 0;
}
```

Here is another version of the above program that does not use the mutex locks: `pthread_sum2.c`